

Solution 4

Solved by Aaron Colin Foote, acf502, 4258770

- Width (total number of blocks at leaf level) of B*-tree implementing an index on *manuf* is equal to w_m .
 w_m is equal to 50.
 - “an attribute *manuf* is indexed”
 - “an attribute *manuf* has 50 distinct values”
- Each block contains 5000/1000 rows.
This means each block records 5 rows (**blocking factor** = 5)
 - “a relational table REPAIR occupies 1000 data blocks”
 - “a relational table REPAIR contains 5000 rows”
- Fanout = 20
 - “all indexes are implemented as B*-trees with a fanout equal to 20”
- $\log_{20} 50 \approx 1.306$
- $\log_{20} 5000 \approx 2.843$

Query 1:

```
SELECT DISTINCT manuf
FROM REPAIR;
```

We have to read w_m blocks.

Index processing:

w_m blocks (horizontal traversal through leaf level on an index on *manuf*)
50 index values read.

Query 2:

```
SELECT rego
FROM REPAIR
WHERE manuf = 'Toyota' AND model = 'Corolla';
```

We have to read $\lceil \log_{20} 50 + 1 \rceil + ((5000/50) + 5000/(50*5))/2$ blocks

Index processing:

$\lceil \log_{20} 50 + 1 \rceil \approx \lceil 1.306 + 1 \rceil \approx 2.306$ blocks

Total number of row identifiers found $5000/50 = 100$ row identifiers

Table processing:

The best case (full clustering): $5000/(50*5) = 20$ blocks

The worst case (each row in a different block): $5000/50 = 100$ blocks

Query 3:

```
SELECT rego
FROM REPAIR
WHERE TO_CHAR(rdate,'YYYY') = '2012';
```

We have to read 1000 data blocks because the entire relational table REPAIR has to be scanned.

Table processing:

1000 blocks (total block capacity of table REPAIR (5000/5 (blocking factor)))

Query 4:

```
SELECT COUNT(*)
FROM REPAIR
WHERE manuf = 'Honda';
```

We have to read $\text{floor}(\log_{20}50 + 1) = 2$ blocks.

Index processing:

$\text{floor}(\log_{20}50 + 1) \approx \text{floor}(1.306 + 1) = 2$ index blocks

Counting row identifiers:

0 data blocks

Query 5:

```
SELECT *
FROM REPAIR
WHERE rego = 'PKR856' AND address = '15 Station St.'
AND rdate = '15-DEC-2010';
```

We have to read $\lceil \log_{20}5000 + 1 \rceil + 1$ blocks because the primary key (rego, address, rdate) is automatically indexed.

$\approx \lceil 2.843 + 1 \rceil + 1$
 ≈ 4.843 blocks

Query 6:

```
SELECT COUNT(model)
FROM REPAIR;
```

We have to read 1000 blocks because the entire relational table REPAIR has to be scanned.

Table processing:

1000 blocks (total block capacity of table REPAIR (5000/5 (blocking factor)))

Query 7:

```
SELECT COUNT(DISTINCT model)
FROM REPAIR;
```

We have to read 1000 blocks because the entire relational table REPAIR has to be scanned.

Table processing:

1000 blocks (total block capacity of table REPAIR (5000/5 (blocking factor)))

Query 8:

```
SELECT manu, COUNT(*)
FROM REPAIR
GROUP BY manu;
```

We have to read w_m ($w_m = 50$) blocks

Index processing:

w_m ($w_m = 50$) blocks

Query 9:

```
SELECT *  
FROM REPAIR  
ORDER BY rego;
```

We have to read 1000 blocks because the entire relational table REPAIR has to be scanned.

Table processing:

1000 blocks (total block capacity of table REPAIR (5000/5 (blocking factor)))

Query 10:

```
SELECT *  
FROM REPAIR  
WHERE rego= 'PKR856' AND manuf = 'Honda';
```

We have to read $\text{floor}(\log_{20} 50 + 1) = 2$ blocks.

Index processing:

$\text{floor}(\log_{20} 50 + 1) \approx \text{floor}(1.306 + 1) = 2$ blocks

Table processing:

The best case (full clustering): $5000/(50 \cdot 5) = 20$ blocks

The worst case (each row in a different block): $5000/50 = 100$ blocks